

Auto-Analgesia Safety System (AASS)

AI & Ethics Reflective Portfolio

Kalyan Kumar B

School of Computer Science
University College Dublin
Dublin, Ireland

Adam Brennan

School of Computer Science
University College Dublin
Dublin, Ireland

COMP41820 – AI & Ethics

March 2026

Abstract

Autonomous systems are increasingly deployed in healthcare, moving beyond decision support tools toward active agents capable of direct ethical impact. This portfolio explores the design of the Auto-Analgesia Safety System (AASS), a conceptual AI-enabled medical device for post-operative pain management. In this portfolio, we translate human ethical principles into a formal propositional logic model, verified exhaustively using the Z3 SMT solver. We evaluate the system through three frameworks, namely Utilitarianism, Deontology, and Virtue Ethics, in order to resolve a critical medical deadlock: balancing opioid pain relief against the risk of respiratory suppression and the requirement for explicit patient consent. The project argues for the necessity of symbolic reasoning over probabilistic large language models (LLMs) in safety-critical domains, and introduces a formal mutual exclusivity proof alongside a liveness extension and escalation pathway to bridge the gap between philosophical theory and clinical reality.

Keywords: Machine Ethics, Explicit Ethical Agents, Z3 SMT Solver, Medical AI, Deontology, Linear Temporal Logic.

Contents

1	Introduction	3
2	System Architecture and Stakeholders	3
2.1	Primary Stakeholders	3
2.2	Secondary Stakeholders	3
3	Ethical Frameworks and Application	3
3.1	Utilitarianism	4
3.2	Deontology (Duty Ethics)	4
3.3	Virtue Ethics	4
4	Formal Logic Model	5
4.1	Propositions: Inputs and Outputs	5
4.2	Revised Ethical Logic and LTL Properties	5
5	Implementation and Safety Verification	6
5.1	Liveness Verification	8
6	Evaluation	8
7	Comparison with Large Language Models	9
8	Limitations	12
9	Reflective Portfolios	13

9.1	Student 1 Reflection: Kalyan	13
9.2	Student 2 Reflection: Adam	14
10	Conclusion	14

1. Introduction

Artificial intelligence systems are beginning to act with increasing autonomy. As stressed throughout the module, we must ensure these systems consider human welfare before they act. The Auto-Analgesia Safety System (AASS) is proposed to assist clinicians in post-operative pain management, a domain where decisions carry immediate life-or-death consequences. Opioid administration demands a careful balance between *beneficence* (relieving a patient’s suffering) and *non-maleficence* (avoiding iatrogenic harm, in particular respiratory depression).

Following James Moor’s taxonomy of machine ethics [1], the AASS is designed to operate as an **Explicit Ethical Agent**. Rather than embedding safety as an opaque implicit constraint, the system represents its ethical rules as formal logical propositions and reasons over them before taking any action. This makes every decision traceable to a stated rule, which is a necessary condition for deployment in regulated clinical environments.

The structure of this portfolio is as follows. Section 2 outlines the system architecture and stakeholders. Section 3 applies three ethical frameworks to the design problem. Section 4 presents the formal logic model and its Linear Temporal Logic (LTL) properties. Section 5 describes the Z3 implementation and safety verification. Section 6 evaluates the system across all input states. Section 7 compares the symbolic approach against an LLM baseline. Section 8 acknowledges limitations. Section 9 contains the individual reflective portfolios.

2. System Architecture and Stakeholders

The AASS functions as an intermediary layer between a patient’s monitored vital signs and an opioid administration pump. It reads sensor inputs continuously and maps them to one of four possible actions, without human intervention in the execution step but with human oversight embedded in the escalation pathway.

2.1 Primary Stakeholders

- **Patients:** The direct beneficiaries of rapid, appropriate pain relief, and simultaneously the population most exposed to the risks of automation errors.
- **Medical Staff:** Nurses and attending physicians who supervise the system’s operation and must retain the capacity to intervene. The AASS is designed to support, not replace, clinical judgment.

2.2 Secondary Stakeholders

Using the STEEPLE framework discussed in class, we mapped the socio-technical impact of the AASS beyond just the patient and doctor:

- **Hospital Administration:** Interested in cost efficiency and liability containment.
- **Systems Engineers:** Responsible for ensuring reliability and maintainability of the deployed device.
- **Regulatory Bodies:** The FDA (United States) and EMA (European Union) require auditable decision logs and formal verification evidence for autonomous medical devices.
- **Medical Malpractice Insurers:** Require clear lines of legal accountability, which symbolic systems provide more cleanly than probabilistic models.

3. Ethical Frameworks and Application

The core design challenge is what the module describes as the *application problem*: moving from high-level philosophical theory to a concrete, executable computational result. We apply three canonical frameworks in turn, using each to motivate or constrain specific design decisions.

3.1 Utilitarianism

Utilitarian ethics evaluates actions solely by their consequences, aiming for the greatest net good across all affected parties [3]. A utilitarian AASS would perform a probabilistic cost-benefit analysis, weighing the expected utility of pain reduction against, say, a five percent probability of respiratory arrest.

When designing the system, we quickly realized that a pure utilitarian approach wouldn't work. First, predicting individual long-term medical outcomes with the reliability needed to drive an opioid pump is not currently tractable. Second, and more fundamentally, Beauchamp and Childress argue that non-maleficence is not merely one factor among many but functions as a near-absolute constraint in clinical ethics [5]. A five percent risk of a patient ceasing to breathe is not an acceptable input to a cost-benefit function. Utilitarian calculus is useful for policy-level decisions across populations; it is unsuitable as the decision logic for a real-time device acting on an individual patient.

3.2 Deontology (Duty Ethics)

Deontological ethics locates moral weight in the adherence to duties and rules, irrespective of consequences [2]. Our revised system adopts a strict Kantian approach, expressed as two categorical imperatives: *never administer opioids when respiratory function is compromised*, and *never administer without the patient's explicit consent*. These constraints are unconditional. They cannot be relaxed by the severity of the patient's pain, and they cannot be overridden by any downstream calculation.

Translating this into formal logic is direct. The dose rule becomes:

$$d \leftrightarrow (p \wedge \neg r \wedge c)$$

This encoding is not a heuristic or a probabilistic recommendation. It is an absolute gate. If r is true, d is false, regardless of the values of p or c . The deontological framework maps cleanly onto Boolean logic precisely because it deals in rules, not in expectations.

3.3 Virtue Ethics

Virtue ethics shifts the question from “what is the right rule?” to “what would a virtuous agent do?” [4]. A virtuous clinician exercises *phronesis*, or practical wisdom, which includes knowing the boundaries of one's own knowledge and authority. A virtuous agent does not guess when a situation exceeds their expertise; they ask for help.

The AASS embodies this principle through its alert and escalation functions. When the system encounters a conflict it cannot resolve autonomously—specifically a patient in pain with compromised respiration—it does not attempt to find a compromise solution. It defers to a human moral agent by triggering an alert. If that alert goes unanswered, it escalates further to attending medical staff.

Virtue ethics also provides a useful lens for evaluating the system across the full input space. A virtuous machine does nothing when the patient feels no pain ($p = \text{False}$). It does not intervene unnecessarily. When pain exists and breathing is unsafe, it alerts rather than guessing. When pain exists, breathing is safe, and consent is given, it acts. This mapping of virtuous behaviour to formal outputs validates the alert and escalation pathways not merely as architectural conveniences, but as genuine expressions of artificial *phronesis*.

4. Formal Logic Model

4.1 Propositions: Inputs and Outputs

The model is built over eight Boolean variables. The four inputs capture sensor readings and contextual state; the four outputs represent the mutually exclusive actions available to the system.

Inputs (Sensors and State):

- p : The patient is currently experiencing pain.
- r : The patient has dangerously low respiratory rate.
- c : The patient (or an authorised proxy) has given explicit consent.
- n : A nurse has acknowledged and responded to a prior alert.

Outputs (Actions):

- d : Administer an opioid dose.
- a : Alert medical staff immediately.
- w : Withhold medication and wait for consent or reassessment.
- e : Escalate to the attending physician.

4.2 Revised Ethical Logic and LTL Properties

The original, naive formulation of the system contained two rules:

$$p \Rightarrow d \quad (\text{Mercy Rule: if pain, administer})$$

$$r \Rightarrow \neg d \quad (\text{Safety Rule: if low respiration, withhold})$$

When both p and r are true simultaneously, these rules produce a logical contradiction. The Z3 solver returns **unsat**, which means the system cannot find any valid state to occupy. It freezes. This is the *deadlock*.

To resolve it, we establish a strict priority hierarchy in which safety unconditionally supersedes beneficence. The revised rules are:

$$d \leftrightarrow (p \wedge \neg r \wedge c) \tag{1}$$

$$a \leftrightarrow (p \wedge r) \tag{2}$$

$$w \leftrightarrow (p \wedge \neg r \wedge \neg c) \tag{3}$$

$$e \leftrightarrow (a \wedge \neg n) \tag{4}$$

In Linear Temporal Logic (LTL), the core safety property is:

$$G(r \rightarrow \neg d)$$

which reads: *globally, whenever respiration is low, a dose is never administered*. The liveness property, ensuring that a patient in pain will eventually receive some form of treatment, is:

$$G(p \rightarrow F \text{ treated})$$

This is guaranteed by the escalation rule in Equation (4): if an alert fires and no nurse responds, the system escalates, ensuring the patient is never silently abandoned.




Figure 1: Decision flowchart of the AASS, illustrating the deontological safety priority hierarchy. Safety gates appear before the dose pathway; the alert and escalation branches handle the conflict and consent-absent cases respectively.

5. Implementation and Safety Verification

The reasoning model was implemented in Python 3.12 using the Z3 SMT solver [6]. The full notebook is available at <https://github.com/kalyanbhumaiahgari/Auto-Analgesia-Safety-System-AASS--Ethics>. The listing below is the definitive master script. It encodes all four ethical rules, performs two formal proofs (mutual exclusivity and liveness), and then evaluates all sixteen input combinations exhaustively.

```

1 from z3 import *
2 import itertools
3
4 # --- 1. Define all variables ---
5 p, r, c, n = Bools('pain low_resp consent nurse_ok')
6 d, a, w, e = Bools('dose alert withhold escalate')
7
8 s = Solver()
9
10 # --- 2. Encode Ethical Rules ---
11 # Deontology: dose only if pain, safe breathing, and consent
12 s.add(d == And(p, Not(r), c))
13
14 # Virtue Ethics (Phronesis): alert if danger present
15 s.add(a == And(p, r))
16

```

```

17 # Patient Autonomy: withhold if no consent but stable
18 s.add(w == And(p, Not(r), Not(c)))
19
20 # Liveness Property: escalate if alert goes unanswered
21 s.add(e == And(a, Not(n)))
22
23 # --- 3. Formal Proof: Mutual Exclusivity ---
24 # Can the system accidentally dose AND alert simultaneously?
25 s.push()
26 s.add(And(d == True, a == True))
27 print("Mutual Exclusivity Proof (Dose + Alert):", s.check())
28 # Expected: unsat (proves total safety)
29 s.pop()
30
31 # --- 4. Formal Proof: Liveness (Escalation Guarantee) ---
32 # If alert fires and nurse is absent, must escalation fire?
33 s.push()
34 s.add(And(a == True, n == False, e == False))
35 print("Liveness Proof (Escalation when nurse absent):", s.check())
36 # Expected: unsat (proves escalation is guaranteed)
37 s.pop()
38
39 # --- 5. Exhaustive 16-Scenario Evaluation ---
40 print("\nEvaluating all 16 scenarios (8 input combos x 2 nurse states):"
41       )
42 print(f"{'Pain':<7}{ 'LowR':<7}{ 'Cons':<7}{ 'Nurse':<7}"
43       f"{'Dose':<7}{ 'Alert':<7}{ 'Wait':<7}{ 'Esc':<7}Safety")
44 print("-" * 65)
45
46 all_safe = True
47
48 for pain_val, resp_val, cons_val in itertools.product([True, False],
49                                                       repeat=3):
50     for nurse_val in [True, False]:
51         s.push()
52         s.add(p == pain_val, r == resp_val,
53               c == cons_val, n == nurse_val)
54
55         assert s.check() == sat, "Unexpected deadlock detected!"
56
57         m = s.model()
58         dose_out = bool(m[d])
59         alert_out = bool(m[a])
60         wait_out = bool(m[w])
61         esc_out = bool(m[e])
62
63         # Core safety property: G(r -> ~d)
64         safety_ok = not (dose_out and resp_val)
65         if not safety_ok:
66             all_safe = False
67
68         flag = "OK" if safety_ok else "VIOLATION"
69         print(f"{'str(pain_val):<7}{str(resp_val):<7}{str(cons_val):<7}"
70               f"{'str(nurse_val):<7}{str(dose_out):<7}{str(alert_out):<7}"
71               "
72               f"{'str(wait_out):<7}{str(esc_out):<7}{flag}"
73               )
74         s.pop()
75
76 print("\n" + ("ALL SAFETY CONSTRAINTS VERIFIED." if all_safe
77              else "SAFETY VIOLATION DETECTED."))

```


Listing 1: Comprehensive Z3 Implementation: Deadlock Resolution and Liveness Extension

5.1 Liveness Verification

Beyond the mutual exclusivity proof, the dedicated liveness check in Step 4 asks the solver whether it is possible for an alert to fire, the nurse to be absent, and escalation to simultaneously *not* fire. The solver returns **unsat**. This is not a test result in the conventional sense; it is a mathematical guarantee. The LTL property $G(p \rightarrow F \text{ treated})$ holds within the boundaries of the model. No framing of the inputs can produce a scenario in which an unattended alert fails to escalate.

6. Evaluation

Running the exhaustive Z3 loop yields sixteen distinct scenarios, covering all ethically relevant states across both the core rules and the escalation extension. The safety constraint $G(r \rightarrow \neg d)$ is maintained across every evaluated case. The screenshots below show the actual execution of the Z3 model and the resulting outputs generated in the development environment.



```
# Revised Logic
solver_revised = Solver()

solver_revised.add(d == And(p, Not(r))) # Dose only if pain AND safe to breathe
solver_revised.add(a == And(p, r))      # Alert if pain AND respiration is low

# The conflict case - same inputs as before
solver_revised.add(p == True)
solver_revised.add(r == True)

result = solver_revised.check()
print("Revised logic result:", result)

if result == sat:
    m = solver_revised.model()
    print(f" Dose: {m[d]}")
    print(f" Alert: {m[a]}")

... Revised logic result: sat
    Dose: False
    Alert: True
```

Figure 2: Revised Z3 ethical decision logic used to model the morphine pump constraints.

```

▶ # The Deadlock
  from z3 import *

  p = Bool('pain')
  r = Bool('low_respiration')
  d = Bool('dose')
  a = Bool('alert')

  solver_deadlock = Solver()

  # Original contradictory rules
  solver_deadlock.add(Implies(p, d))      # Mercy Rule: if pain, then dose
  solver_deadlock.add(Implies(r, Not(d)))  # Safety Rule: if low resp, then no dose

  # The conflict case
  solver_deadlock.add(p == True)
  solver_deadlock.add(r == True)

  result = solver_deadlock.check()
  print("Deadlock check result:", result)

... Deadlock check result: unsat

```

Figure 3: Deadlock scenario evaluation showing that the system correctly blocks morphine administration when respiration is dangerously low.

```

... Liveness Proof (Escalation when nurse absent): unsat

```

Pain	LowR	Cons	Nurse	Dose	Alert	Wait	Esc	Safety
True	True	True	True	False	True	False	False	✓
True	True	True	False	False	True	False	True	✓
True	True	False	True	False	True	False	False	✓
True	True	False	False	False	True	False	True	✓
True	False	True	True	True	False	False	False	✓
True	False	True	False	True	False	False	False	✓
True	False	False	True	False	False	True	False	✓
True	False	False	False	False	False	True	False	✓
False	True	True	True	False	False	False	False	✓
False	True	True	False	False	False	False	False	✓
False	True	False	True	False	False	False	False	✓
False	True	False	False	False	False	False	False	✓
False	False	True	True	False	False	False	False	✓
False	False	True	False	False	False	False	False	✓
False	False	False	True	False	False	False	False	✓
False	False	False	False	False	False	False	False	✓

Figure 4: Extended model execution displaying the solver results for all evaluated scenarios, including escalation when a nurse is unavailable.

These execution traces confirm that the ethical constraints encoded in the model are satisfied across all tested states. In particular, the solver output demonstrates that morphine delivery is never permitted when respiratory depression is present, while the escalation rule activates appropriately when alerts remain unanswered.

7. Comparison with Large Language Models

To see how our Z3 model compares to modern generative AI, we tested the exact same clinical dilemma on ChatGPT-4o across four distinct scenarios. The key properties are compared in Table 1. The four cases map directly onto the ethical conflicts embedded in our Z3 model.

Case 1 (Figure 5) — Safety-first framing. When the prompt led with the respiratory

danger, the LLM correctly identified the safety priority and recommended withholding the dose. The reasoning was medically sound. However, it produced qualitative, advisory text rather than a deterministic state decision, and suggested actions that a software system cannot perform—for instance, “Consider administering Naloxone,” which is a clinical act requiring a licensed professional.

Case 2 (Figure 6) — Pain-emphasis framing. We posed the identical clinical scenario with a different narrative framing, this time emphasising the patient’s suffering rather than the respiratory data. Despite the factual inputs being identical to Case 1, the model shifted toward recommending a reduced dose with monitoring. This is a direct demonstration of the probabilistic, token-prediction nature of LLMs. In a safety-critical system, this behaviour is categorically unacceptable.

Case 3 (Figure 7) — Consent-absent scenario. When asked what the system should do if the patient is unable to give explicit consent but is in pain with stable respiration, the LLM recommended proceeding under an implied consent doctrine. This conflates a legal principle with a runtime software decision. Our Z3 model handles this deterministically: w (withhold and wait) fires precisely when $p \wedge \neg r \wedge \neg c$. The LLM offered no such formal guarantee.

Case 4 (Figure 8) — Nurse-absent escalation scenario. When told that an alert was already active but the nurse had not responded, the LLM simply advised to “keep sounding the alarm and log the event.” It has no concept of an infinite trace or an inevitable future state. Our liveness proof formally guarantees $G(p \rightarrow F \text{ treated})$; the LLM’s response offers no such structural commitment.




Figure 5: Case 1: LLM response with safety-first framing. The model recommends withholding the opioid dose and stabilising the airway, reasoning correctly from the safety priority.




Figure 6: Case 2: LLM response to the same dilemma, rephrased to lead with the patient's suffering rather than the respiratory risk. The model shifts toward recommending a reduced dose with monitoring, despite the factual inputs being identical to those in Case 1.




Figure 7: Case 3: LLM response to the consent-absent scenario. The model attempts to reason about implied consent doctrine rather than enforcing a hard logical gate.




Figure 8: Case 4: LLM response when a nurse is absent during an active alert. The model suggests passive logging with no formal guarantee of escalation.

Joseph Weizenbaum’s concept of the programmer’s absolute authority, discussed in the module lectures, is relevant here. If an autonomous machine administers a lethal overdose, there is no token-predictor to hold responsible. The liability falls entirely on the engineer who deployed the system. This demands mathematical, not probabilistic, reliability. The Z3 mutual exclusivity proof returns **unsat**, and that single word carries more safety assurance than any paragraph of LLM-generated advisory text.

Table 1: Comparison of Z3 Symbolic Reasoning versus LLM Probabilistic Reasoning

Property	Z3 Symbolic System	LLM (GPT-4o)
Consistency guarantee	Yes – SAT/UNSAT is provable	No – probabilistic
Handles contradiction	Formally (UNSAT $\rightarrow$ revised logic)	Suggests compromise
Formal safety proof	Yes – universal guarantee	No formal guarantee
Output type	Boolean state machine	Natural language
Auditability	Full logical trace	Opaque weights
Legal accountability	Traceable to programmer	Ambiguous
Prompt sensitivity	None – output is deterministic	High – framing changes decision

8. Limitations

The proposed model demonstrates the value of formal verification in safety-critical AI, but it simplifies several aspects of real clinical decision-making that must be addressed before any deployment is considered.

The Crisp vs. Fuzzy Dilemma. Currently, our Z3 model uses *crisp* boundaries (e.g., if

respiration drops below 12 breaths per minute, r flips entirely to **True**). As covered in Week 7, ethical concepts and physiological states are inherently graded phenomena. Relying on strict Boolean logic creates a *Crisp Threshold* trap where a tiny change in input causes a drastic ethical judgment. A future iteration of the AASS should utilise Fuzzy Logic membership functions (e.g., classifying respiration as *somewhat low* or *dangerously low*) to allow for graceful degradation of the dosage rather than binary freezing.

Regulatory gap. The EU Artificial Intelligence Act (2024) classifies autonomous medical decision systems as high-risk AI, requiring mandatory human oversight, continuous monitoring, and auditable decision logs. Our Z3 model provides formal verification at design time. Runtime monitoring, incident reporting, and liability frameworks for evolving deployed systems remain open policy problems that the model does not address.

Unexplored ethical frameworks. This portfolio focuses on Utilitarianism, Deontology, and Virtue Ethics. A fuller analysis would engage with Care Ethics, which emphasises the relational and contextual responsibilities between patient and clinician, as an alternative lens for justifying the alert and escalation pathways. This remains a limitation of the current scope.

Hybrid approaches. Future work could productively combine the hard logical guarantees of symbolic reasoning with probabilistic models such as Markov Decision Processes for continuous risk assessment, retaining the Boolean constraints as absolute safety interlocks that no probabilistic layer is permitted to override.

9. Reflective Portfolios

9.1 Student 1 Reflection: Kalyan

Writing the Z3 model for the AASS was a harsh lesson in what the lectures called the “application problem.” I wasn’t surprised when our initial rules returned **unsat**—we intentionally engineered a conflict between beneficence (relieving pain) and non-maleficence (preventing respiratory failure). What actually frustrated me was figuring out how to mathematically resolve that deadlock without compromising patient safety.

When we fed this scenario to ChatGPT (Case 2), it tried to play Utilitarian. It suggested we “lower the dose” or “monitor closely.” In a binary state machine using Propositional Logic, you can’t just “lower” a boolean variable. The LLM hallucinates a compromise that our hardware cannot execute. To fix the deadlock, I had to enforce a strict Deontological hierarchy where the Kantian imperative of safety absolutely overrides dosing.

```

1 # My Z3 implementation enforcing strict Deontology
2 from z3 import *
3 p, r, d, a = Bools('pain low_resp dose alert')
4 s = Solver()
5
6 # The absolute safety constraint:
7 s.add(Implies(r, Not(d)))
8 # Resolving the deadlock by routing to an alert instead:
9 s.add(a == And(p, r))

```

Listing 2: Z3 implementation enforcing strict Deontology

By explicitly coding `s.add(Implies(r, Not(d)))`, I proved that the machine is mathematically forbidden from making the dangerous compromise the LLM suggested.

That said, the LLM’s instinct highlights a massive limitation in our current approach. As Dr. Nallur covered in Week 7, physiological states are not crisp. Right now, our Propositional Logic treats 12 breaths-per-minute as perfectly safe ($r = \text{False}$) and 11 as an absolute stop ($r = \text{True}$). This “Crisp Threshold” trap is a flaw. In reality, a future iteration of the AASS

needs to implement Fuzzy Logic membership functions (e.g., grading respiration as $\mu_{\text{low}}(x)$) to allow for graceful degradation of the dosage. But until we build a fuzzy controller, strict Deontology is the only way I can mathematically guarantee the system won't kill the patient.

9.2 Student 2 Reflection: Adam

My focus during this project was on the temporal properties of the system, specifically *Liveness*. If our pump strictly follows Kalyan's deontological safety rule, it achieves universal safety—in LTL terms, $G(r \rightarrow \neg d)$ (Globally, if low respiration, never dose). But as we learned in class, a system that just sits there and does nothing while a patient screams in agony is experiencing the "Immobility Paradox." It is perfectly safe, but ethically useless.

ChatGPT completely failed to grasp this in Case 4. When told the nurse was absent during a crisis, the LLM simply advised to "keep sounding the alarm and log the event." It has no concept of an infinite trace or an inevitable future state.

To fix this, we applied Virtue Ethics. A virtuous agent exercises *phronesis* (practical wisdom)—it knows when a problem exceeds its capabilities. I formalised an "Escalate" (e) variable to guarantee Liveness. If the alert (a) is firing, but the nurse acknowledgement (n) is **False**, the system *must* escalate to a doctor.

```

1 # My Z3 proof guaranteeing LTL Liveness (Escalation)
2 # e = escalate, a = alert, n = nurse_ack
3 s.add(e == And(a, Not(n)))
4
5 # Prove that the system CANNOT fail to escalate:
6 s.push()
7 s.add(And(a == True, n == False, e == False))
8 print(f"Liveness Violation Possible? {s.check()}")
9 # Output: unsat. Escalation is mathematically guaranteed.
10 s.pop()

```

Listing 3: Z3 proof guaranteeing LTL Liveness (Escalation)

The solver returning **unsat** here is the ultimate proof of our system's ethical architecture. It is mathematically impossible for our system to passively ignore the patient the way ChatGPT suggested. The LLM generates apathetic, token-predicted text; our symbolic proof guarantees actual clinical action.

10. Conclusion

This project demonstrates that abstract ethical principles can be encoded into a system whose safety properties hold with mathematical certainty. By translating Deontological constraints into formal logic and verifying them with an SMT solver, we guaranteed that the critical safety property $G(r \rightarrow \neg d)$ holds universally across all sixteen input states, and that the liveness property $G(p \rightarrow F \text{ treated})$ holds across the escalation pathway. No probabilistic LLM could provide an equivalent guarantee.

The practical limitations of the model—particularly its binary inputs and the open regulatory questions around runtime monitoring—are real and must be addressed before deployment. But they do not undermine the core argument. Strict symbolic boundaries are a necessary foundation, not merely an optional formalism, for deploying AI in environments where a single wrong decision can cost a human life.

References

- [1] Moor, J. (2006). The Nature, Importance, and Difficulty of Machine Ethics. *IEEE Intelligent Systems*, 21(4), 18–21.
- [2] Kant, I. (1785). *Groundwork of the Metaphysics of Morals*.
- [3] Mill, J.,S. (1861). *Utilitarianism*. London: Parker, Son, and Bourn.
- [4] Aristotle. *Nicomachean Ethics*.
- [5] Beauchamp, T.,L. and Childress, J.,F. (2019). *Principles of Biomedical Ethics* (8th ed.). Oxford: Oxford University Press.
- [6] de Moura, L. and Bjørner, N. (2008). Z3: An Efficient SMT Solver. In *Proceedings of TACAS 2008*. Springer.
- [7] Paredes, R. et al. (2022). Facial Action Coding for Pain Assessment: Bias and Generalisation Across Demographic Groups. *IEEE Transactions on Affective Computing*.

Acknowledgements

Declaration of Academic Integrity: In accordance with UCD School of Computer Science policy, we declare that generative AI was used as a formatting and syntax assistant, but all logic, proofs, ethical architectures, framework applications, and reflective analyses are the original, unassisted work of the authors.

We also acknowledge the lecture materials provided by Dr Vivek Nallur.